


# Network Generation by Deep Reinforcement Learning

Yang Yang, Jiarong Xu  
Zhejiang University

# Why Networks?


- Networks are a natural language for describing and modeling complex systems.



Independent Data


# Why Networks?

- Networks are a natural language for describing and modeling complex systems.




Network!

# Application: Vertex Classification



# Noisy Network Data

- In real world, network data is **hard to obtain** and consists of **noises**
  - Noises are caused by **incomplete and biased data sampling**, **human subjectivity**, and **inconsistent with the downstream task**.

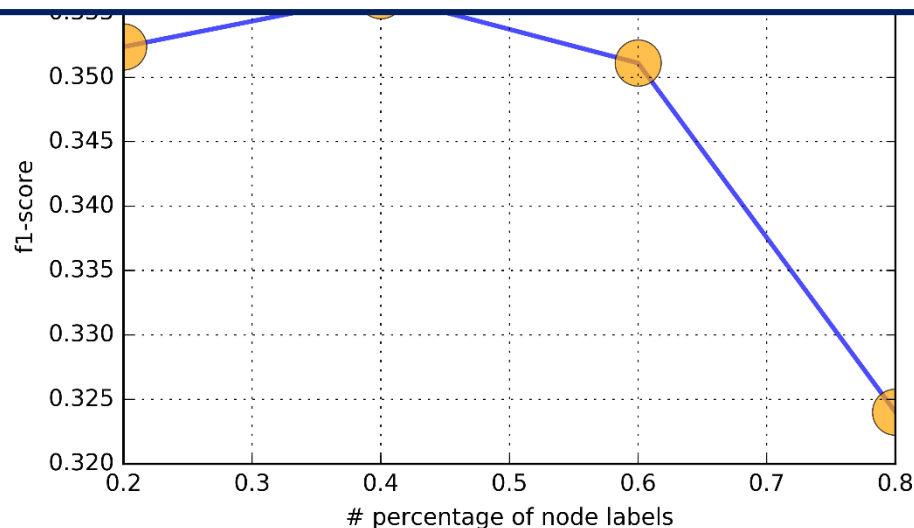


**Task performance drops as more node labels are obtained!**

# Noisy Network Data

- In real world, network data is **hard to obtain** and consists of **noises**
  - Noises are caused by **incomplete and biased data sampling**, **human subjectivity**, and **inconsistent**

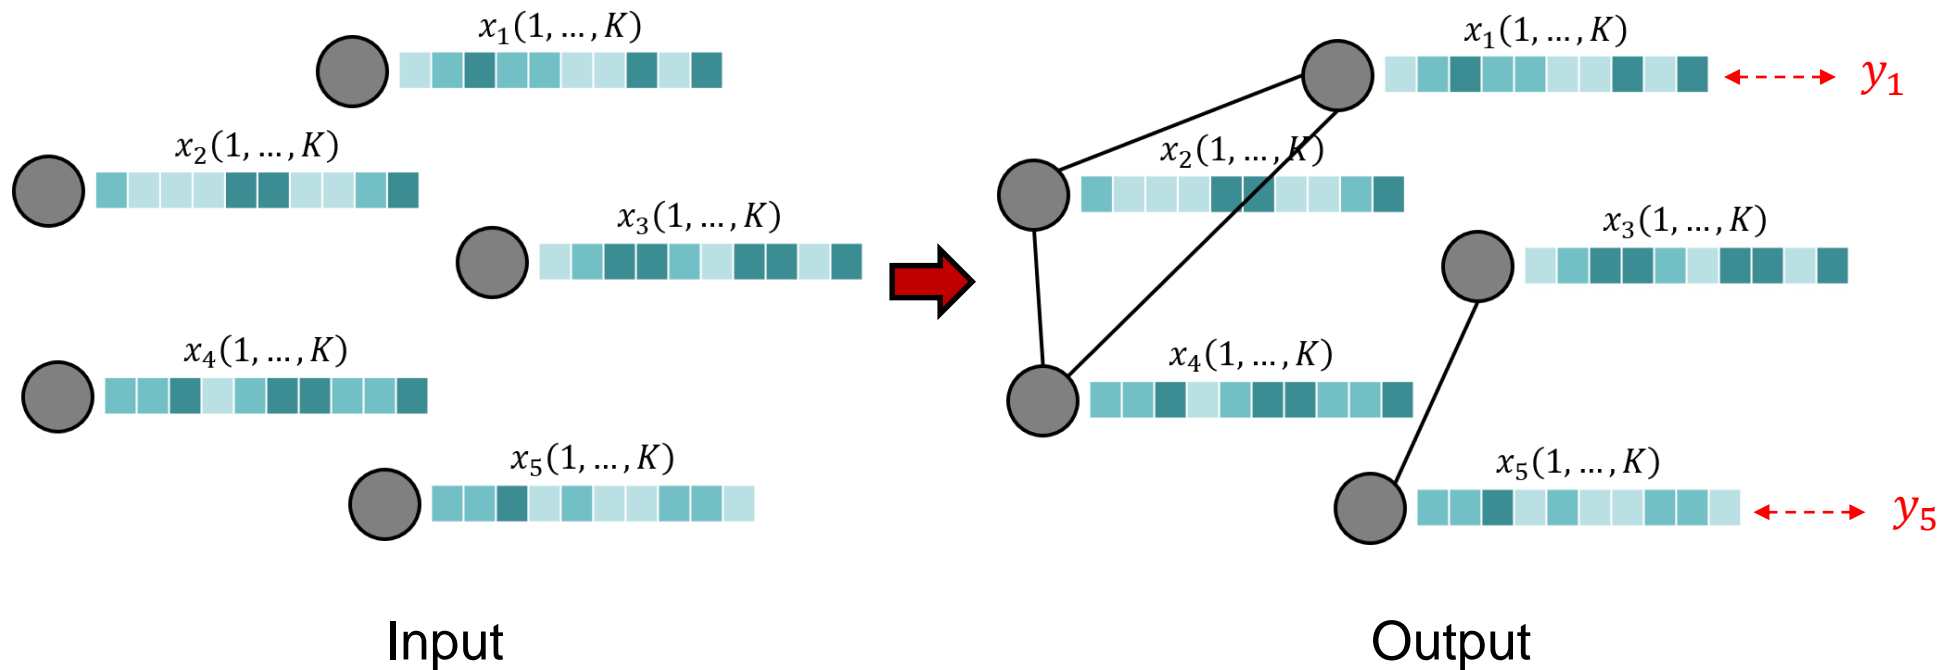
**How to construct a reliable network?**



**Task performance drops as more node labels are obtained!**

# Task

- Construct network by considering **node features** and optimizing **task performance**.




# Related Work

- **Link prediction: given two nodes, determine if there exist an edge.**
  - Hard to preserve macro-level network properties.
- **Network generation model: design a certain mechanism to generate (nodes and) edges.**
  - Requires rich domain knowledge and hard to generalize.



# General Idea

- A network is consisted by several **paths** among nodes.



Path generation can be formulated as Markov Decision Process!

# Our Approach

- Reinforcement learning framework
  - State:**  $(v_0, \dots, v_t)$ , where  $v_t$  is the current node
  - Action:**  $v_{t+1}$ , create an edge between  $v_t$  and  $v_{t+1}$

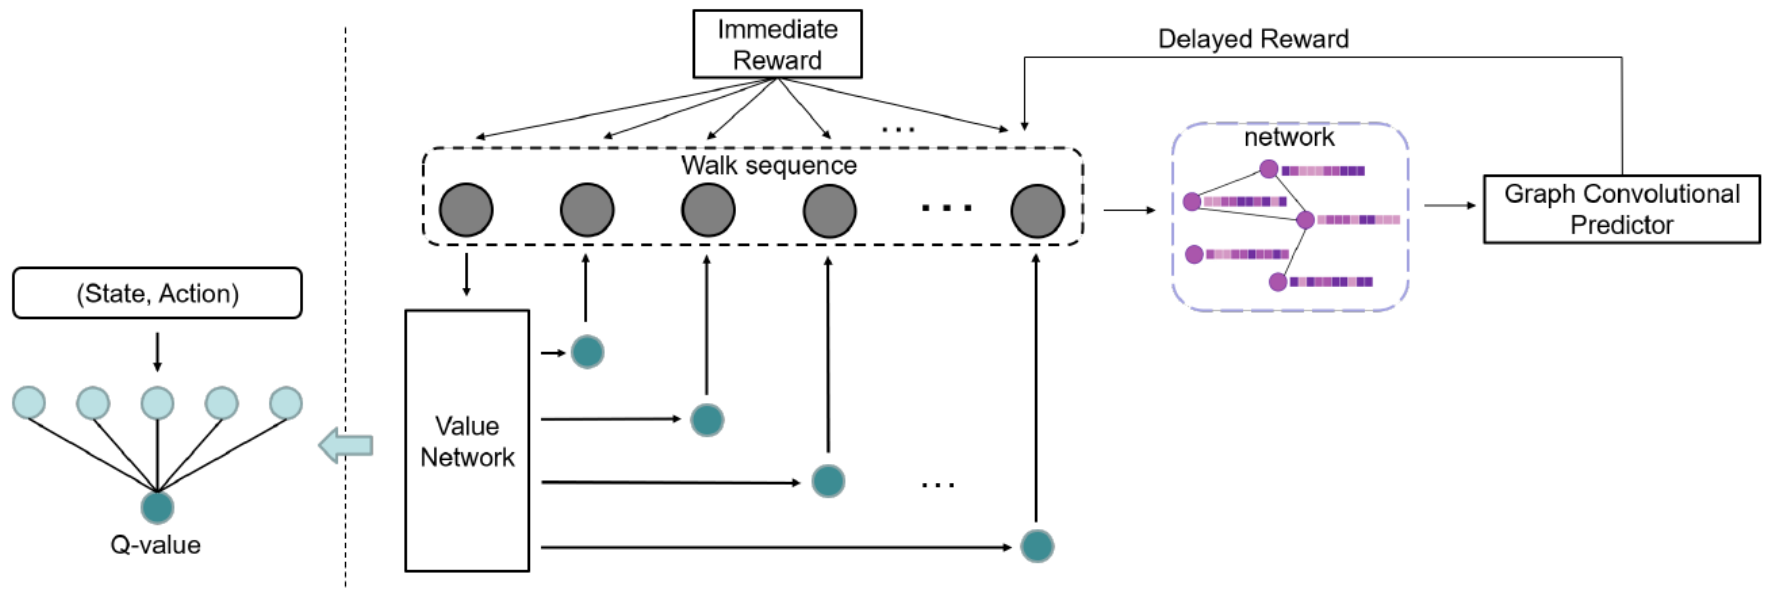


Figure 1: An overview of the proposed network reconstruction method.

# Our Approach

- Reinforcement learning framework

- **Reward:**

- **Immediate reward:** a given network (as train data) guided

$$r_i(v_{t+1}|v_t) = \begin{cases} 0, & (v_t, v_{t+1}) \notin E \text{ or } (v_t, v_{t+1}) \in ((v_0, v_1), (v_1, v_2), \dots, (v_{t-1}, v_t)) \\ 1, & (v_t, v_{t+1}) \in E \end{cases}$$

edges set

- **Delayed reward:** task guided

$$r_d(v_T|(v_0, v_1, \dots, v_T)) = GCN(G_{net}((v_0, v_1, \dots, v_T)), features)$$

# Learning Algorithm

---

**Algorithm 2** Overall Training Process

---

Initialize positive replay memory  $\mathcal{D}_p$  to capacity  $N$   
Initialize negative replay memory  $\mathcal{D}_n$  to capacity  $M$   
Initialize action-value function  $Q$  with random weights

**repeat**

Initialize sequence  $s_0 = (v_0)$

**for**  $t = 0$  **to**  $T$  **do**

With probability  $\varepsilon$  select a random action  $a_t$

Otherwise select  $a_t = \max_a Q^*(v_t, a; \theta)$

Execute action  $a_t$  in emulator and observe reward  $r_t$  and next node  $v_{t+1}$

Set  $s_{t+1} = (v_0, \dots, v_{t+1})$

**if**  $r_t > 0$  **then**

Store transition  $(s_t, a_t, r_t, s_{t+1})$  in  $\mathcal{D}_p$

**else**

Store transition  $(s_t, a_t, r_t, s_{t+1})$  in  $\mathcal{D}_n$

**end if**

Sample random minibatch of transitions  $(s_j, a_j, r_j, s_{j+1})$  from  $\mathcal{D}_p$  and  $\mathcal{D}_n$  with a proportion of 1:4

Set  $y_j = \begin{cases} \text{GCN}(G\_net(s_{j+1}), \text{features}) & \text{len}(s_{j+1}) == T \\ r_j + \gamma \max_{a'} Q(s_{j+1}, a'; \theta) & \text{others} \end{cases}$

Perform a gradient descent step on  $(y_j - Q(s_j, a_j; \theta))^2$

**end for**

**until** episodes finish

---

Select & execute actions with e-greedy

Construct experience replay

Train Q network

# Experimental Setup

- Dataset:
  - Terrorist network:
    - Terrorists(1206) vs. Non-terrorists(2767)
    - Each node has 45 features
  - Citation network: Cora & Pubmed
    - Vertices: documents
    - Edges: citation links between documents
    - Features: sparse bag-of-words representation

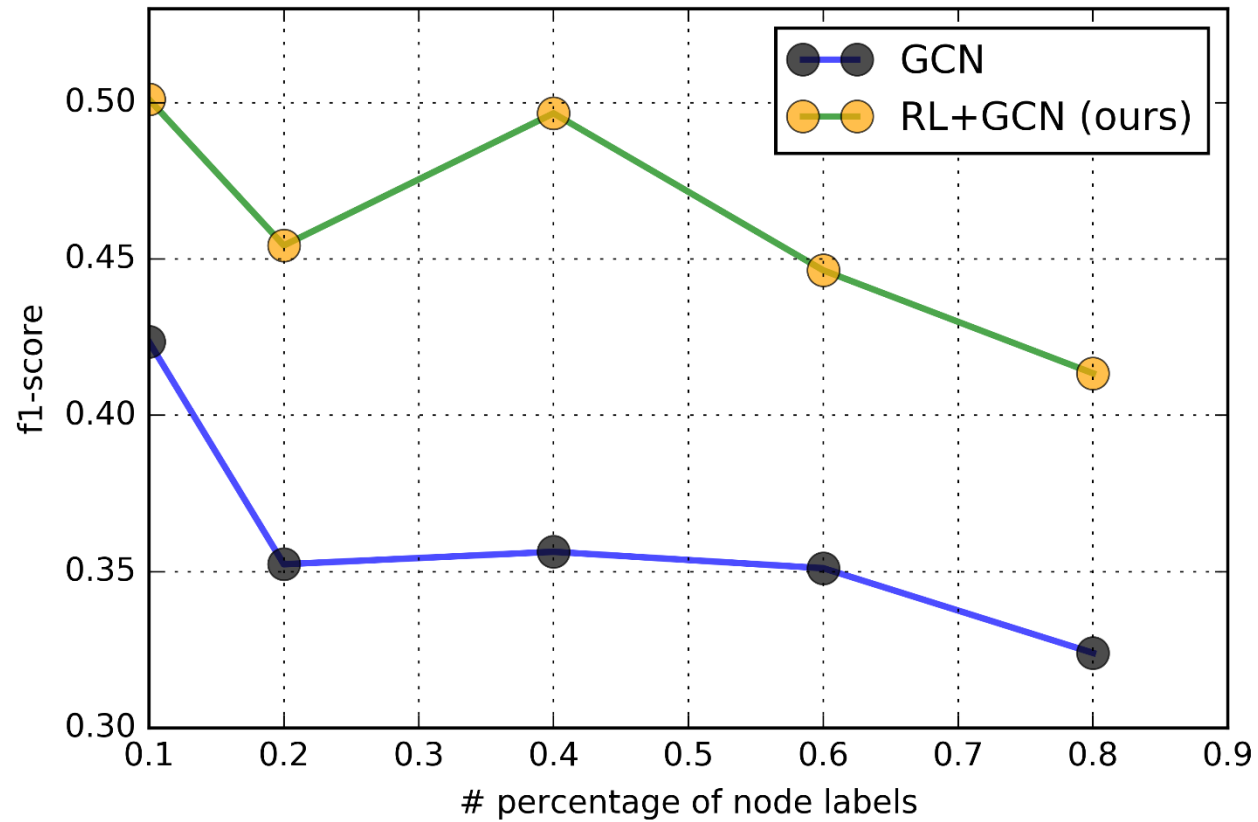
DATASET	TYPE	#VERTICES	#EDGES	#CLASSES	#FEATURES
Hikvision	Terrorist Network	3,973	7,600	2	45
Cora	Citation network	2,708	5,429	7	1,433
Pubmed	Citation network	19,717	44,338	3	500

- Downstream task: **node classification**


# Experimental Setup

- Baselines:
  - **MLP**: Conduct MLP classification based on features
  - **LP**: Construct similarity matrix based on features, and conduct label propagation on it.
  - **RL + LP**: Extract similarity matrix from RL, and conduct label propagation on it.
  - **GCN**: Conduct GCN classification based on pre-built network and features.
  - **RL + GCN**: Reconstruct network from RL, and conduct GCN classification based on the reconstructed network and features.

# RL Does Help!



# Experimental Results – Terrorist Network



With the percentage of node labels **decreases**:

- Method based on **features** gets **worse**
- Method based on **both features and network** gets **better**



# Experimental Results

			MLP	LP	RL+LP	GCN	RL+GCN
# percentage of node labels	0.8	Accuracy	0.82087	0.75536	0.37872	0.69560	0.42893
		Precision	0.73687	0.70622	0.28969	0.39456	0.42893
		Recall	0.56500	0.32961	0.28969	0.27488	0.75829
		F1-score	0.63910	0.44826	0.35928	0.32402	0.41344
	0.6	Accuracy	0.81544	0.73868	0.29811	0.68616	0.43837
		Precision	0.71940	0.67969	0.29198	0.40663	0.30612
		Recall	0.55546	0.31843	0.87212	0.30892	0.82380
		F1-score	0.62651	0.43276	0.43748	0.35111	0.44637
	0.4	Accuracy	0.81112	0.72945	0.32257	0.68330	0.59941
		Precision	0.70925	0.70667	0.32257	0.42828	0.38861
		Recall	0.55789	0.27568	1.00000	0.30511	0.68759
		F1-score	0.62407	0.39663	0.48779	0.35635	0.49657
	0.2	Accuracy	0.80441	0.69991	0.31331	0.66467	0.39321
		Precision	0.69573	0.54907	0.31331	0.40903	0.30908
		Recall	0.55748	0.23594	1.00000	0.30950	0.85699
		F1-score	0.61838	0.33006	0.47713	0.35237	0.45432
	0.1	Accuracy	0.80033	0.70274	0.30425	0.66303	0.52657
		Precision	0.68978	0.52515	0.30425	0.41698	0.36654
		Recall	0.55646	0.23989	1.00000	0.30633	0.79143
		F1-score	0.61545	0.32934	0.46655	0.35319	0.50103

# Experimental Results – Terrorist Network

			MLP	LP	RL+LP	GCN	RL+GCN
# percentage of node labels	0.8	Accuracy	0.82087	0.75536	0.37872	0.69560	0.42893
		Precision	0.73687	0.70622	0.28969	0.39456	0.42893
		Recall	0.56500	0.32961	0.28969	0.27488	0.75829
		F1-score	0.63910	0.44826	0.35928	0.32402	0.41344
	0.6	Accuracy	0.81544	0.73868	0.29811	0.68616	0.43837
		Precision	0.71940	0.67969	0.29198	0.40663	0.30612
		Recall	0.55546	0.31843	0.87212	0.30892	0.82380
		F1-score	0.62651	0.43276	0.43748	0.35111	0.44637
	0.4	Accuracy	0.81112	0.72945	0.32257	0.68330	0.59941
		Precision	0.70925	0.70667	0.32257	0.42828	0.38861
		Recall	0.55789	0.27568	1.00000	0.30511	0.68759
		F1-score	0.62407	0.39663	0.48779	0.35635	0.49657
	0.2	Accuracy	0.80441	0.69991	0.31331	0.66467	0.39321
		Precision	0.69573	0.54907	0.31331	0.40903	0.30908
		Recall	0.55748	0.23594	1.00000	0.30950	0.85699
		F1-score	0.61838	0.33006	0.47713	0.35237	0.45432
	0.1	Accuracy	0.80033	0.70274	0.30425	0.66303	0.52657
		Precision	0.68978	0.52515	0.30425	0.41698	0.36654
		Recall	0.55646	0.23989	1.00000	0.30633	0.79143
		F1-score	0.61545	0.32934	0.46655	0.35319	0.50103

However, the network information hurts the performance overall...  
Why?

# Labels Sensitivity

Table 3: Classification performance on noisy citation network (add 60% noises) with different ratios of vertex labels

		80%		60%		40%	
		micro-f1	macro-f1	micro-f1	macro-f1	micro-f1	macro-f1
Noisy Cora	LP	0.16236	0.05799	0.12177	0.04052	0.12308	0.04322
	RL+LP	0.42620	0.42620	0.13284	0.03350	0.14277	0.05799
	GCN	0.61808	0.48967	0.60526	0.48245	0.60283	0.48175
	<b>RL+GCN (ours)</b>	<b>0.76384</b>	<b>0.71929</b>	<b>0.73155</b>	<b>0.67555</b>	<b>0.72246</b>	<b>0.67023</b>
Noisy Pubmed	LP	0.34559	0.32646	0.38177	0.27309	0.39887	0.29663
	RL+LP	0.38565	0.18554	0.39318	0.20926	0.39363	0.24846
	GCN	0.71661	0.70437	0.71991	0.70622	0.72292	0.70830
	<b>RL+GCN (ours)</b>	<b>0.86105</b>	<b>0.86086</b>	<b>0.86332</b>	<b>0.86308</b>	<b>0.85775</b>	<b>0.85739</b>
		20%		10%			
		micro-f1	macro-f1	micro-f1	macro-f1		
Noisy Cora	LP	0.15136	0.03997	0.15135	0.04649		
	RL+LP	0.12921	0.03269	0.13002	0.03288		
	GCN	0.61343	0.54758	0.60008	0.56347		
	<b>RL+GCN (ours)</b>	<b>0.70051</b>	<b>0.65302</b>	<b>0.66571</b>	<b>0.62909</b>		
Noisy Pubmed	LP	0.39952	0.19033	0.35405	0.24257		
	RL+LP	0.39565	0.26209	0.38905	0.18680		
	GCN	0.72291	0.71060	0.70830	0.70830		
	<b>RL+GCN (ours)</b>	<b>0.85280</b>	<b>0.85251</b>	<b>0.84419</b>	<b>0.84352</b>		

# Noises Sensitivity

Table 4: Classification performance on citation network (given 80% vertex labels) with different ratios of noises

		20% noises		40% noises		60% noises		80% noises	
		micro-f1	macro-f1	micro-f1	macro-f1	micro-f1	macro-f1	micro-f1	macro-f1
Cora	GCN	<b>0.79612</b>	<b>0.74578</b>	0.71089	0.61443	0.61808	0.48967	0.47343	0.33818
	<b>RL+GCN (ours)</b>	0.78044	0.744	<b>0.76753</b>	<b>0.71932</b>	<b>0.76384</b>	<b>0.71929</b>	<b>0.78044</b>	<b>0.74580</b>
Pubmed	GCN	0.71660	0.70437	0.77350	0.76704	0.71661	0.70437	0.67135	0.65223
	<b>RL+GCN (ours)</b>	<b>0.86055</b>	<b>0.86031</b>	<b>0.85928</b>	<b>0.85959</b>	<b>0.85928</b>	<b>0.85941</b>	<b>0.85953</b>	<b>0.85972</b>

# Summary

- We study the problem of **network reconstruction** in a **deep reinforcement learning framework**.
- Our method efficiently improve the downstream task performance, comparing with other method utilizing network information.
- Issues remain:
  - MLP only based on features performs best?
  - Do the generated edges have a clear physical meaning?
  - **We need to dig deeper on the data!**

Thank you!  
Q&A